

Texture Mapping

Prof. Vladlen Koltun
Computer Science Department
Stanford University


How do we simulate surface detail?




- Store functions in memory buffers and map them onto surfaces
- Functions can be generated through simulation or acquisition from the real world

Enhanced realism

Not just bitmaps



one-dimensional



two-dimensional



three-dimensional

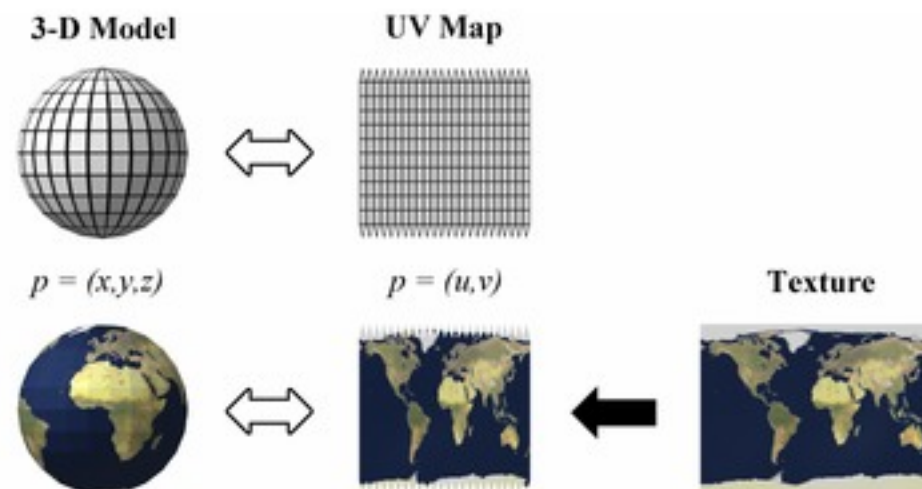
Not really texture



≠




UV mapping




- Associate texture coordinates (u,v) with vertices
- Interpolate within polygons
- Use texture values to determine reflection coefficients

Perspective-correct interpolation essential




linear




perspective-correct

Aliasing




unfiltered





filtered



Texture filtering



Minification and magnification




Bilinear filtering





- Instead of rounding interpolated (u,v) coordinates to integer values, blend the four adjacent texels. Weigh contributions based on distance.
- Aimed to address aliasing artifacts due to magnification

Mipmapping




- Pre-filter the texture map at multiple resolutions, store the hierarchy
- “mip” stands for “multum in parvo”, meaning “much in little”
- Aimed to address aliasing artifacts due to minification

Mipmapping



- 1/3 increase in memory requirements
- Key simplification: approximate the pre-image of a pixel by a square


Mipmapping




$$D = \max \left(\sqrt{\left(\frac{\partial u}{\partial x} \right)^2 + \left(\frac{\partial v}{\partial x} \right)^2}, \sqrt{\left(\frac{\partial u}{\partial y} \right)^2 + \left(\frac{\partial v}{\partial y} \right)^2} \right) \quad d = \log D$$

- Hierarchy indexed by mipmap level d , which ranges from 0 to $\log(w)$
- Texture value produced by bilinear interpolation within two mipmap levels adjacent to d , then linear interpolation between the two values


Mipmapping



Anisotropic Filtering




Summed-Area Tables




- Precompute an “image” S in which every pixel (x,y) stores the sum of values of all pixels (i,j) in the texture such that $i \leq x$ and $j \leq y$.
- For each fragment, compute the texture coordinates for the lower left corner and the upper right corner of its pixel: (u_{ll}, v_{ll}) , (u_{ur}, v_{ur})

Summed-Area Tables



- Precompute an “image” S in which every pixel (x,y) stores the sum of values of all pixels (i,j) in the texture such that $i \leq x$ and $j \leq y$.
- For each fragment, compute the texture coordinates for the bottom left corner and the upper right corner of its pixel: $(u_{ll}, v_{ll}), (u_{ur}, v_{ur})$
- Take
$$\frac{S(u_{ur}, v_{ur}) - S(u_{ll}, v_{ur}) - S(u_{ur}, v_{ll}) + S(u_{ll}, v_{ll})}{(u_{ur} - u_{ll})(v_{ur} - v_{ll})}$$

Even More Accurate Filtering



$$D = \min \left(\sqrt{\left(\frac{\partial u}{\partial x}\right)^2 + \left(\frac{\partial v}{\partial x}\right)^2}, \sqrt{\left(\frac{\partial u}{\partial y}\right)^2 + \left(\frac{\partial v}{\partial y}\right)^2} \right) \quad d = \log D$$

- Take the mipmap level determined by the shorter edge of the pre-image
- Rasterize the pre-image as line segment
- Slower, but advantageous for narrow elongated pre-images